
Training Dense Object Nets: A Novel Approach

Anonymous Author(s)

Affiliation
Address
email

Abstract

1 Our work proposes a novel framework that addresses the computational limitations
2 associated with training Dense Object Nets (DON) while achieving robust and
3 dense visual object descriptors. DON's descriptors are known for their robustness
4 to viewpoint and configuration changes, but training these requires image pairs
5 with computationally expensive correspondence mapping. This limitation hampers
6 dimensionality and robustness, thereby restricting object generalization. To over-
7 come this, we introduce data generation procedure using synthetic augmentation
8 and a novel deep learning architecture that produces denser visual descriptors
9 with reduced computational demands. Notably, our framework eliminates the
10 need for image pair correspondence mapping and showcases its application in a
11 robotic grasping pipeline. Experimental results demonstrate that our approach
12 yields descriptors as robust as those generated by DON.

13

1 Introduction

14 Creating a general-purpose robotic system capable of performing practical tasks, such as those
15 portrayed in movies like Chappie [1] or C-3PO [2], is a significant challenge in robotics. While
16 advancements have been made in related domains, achieving this goal remains an ongoing endeavour.
17 Recent artificial intelligence (AI) breakthroughs, particularly in deep learning, have demonstrated
18 remarkable capabilities. For instance, AlphaGo [3], an AI system trained entirely on self-play,
19 defeated the world's best human Go player at the time. Subsequent algorithms, such as those
20 developed by Silver et al. [4], have mastered complex games like chess, Go, World of Warcraft [5],
21 and Shogi, surpassing human expertise. These achievements underscore the importance of visual
22 data in deep learning, as these algorithms learn directly from visual inputs like gameplay recordings
23 or online video streams. Additionally, the introduction of AlexNet [6] in 2012 has revolutionized
24 computer vision, leading to significant progress in tasks like semantic segmentation [7], object
25 identification and recognition [8], and human pose estimation [9].

26 Significant strides have been made in robotics, including self-driving cars and humanoid robots
27 with advanced capabilities. Integration of AI models like ChatGPT [10] and PaLM [11] enhances
28 human-robot interactions and object perception. However, deploying Large Language Models
29 (LLMs) such as ChatGPT and Palm-E poses challenges. Their size and complexity require substantial
30 computational resources, raising concerns about energy consumption, environmental impact, and
31 the digital divide [12, 13]. Responsible resource allocation and addressing ethical implications are
32 crucial for balancing progress and sustainability.

33 Currently, in robotics, typical industrial robots perform repetitive operations based on pre-programmed
34 instructions, finding the ideal object representation for grasping and manipulation tasks still needs
35 to be answered. Existing representations may be unable to understand an object's geometrical
36 and structural information, rendering them unsuitable for complex tasks. In recent work, Florence
37 et al. [14] introduced a novel visual object representation termed "dense visual object descriptors" to
38 the robotics community. This representation, generated by the Dense Object Nets (DON) framework,

39 converts each pixel in an image ($I[u, v] \in \mathbb{R}^3$) into a higher-dimensional embedding ($I_D[u, v] \in \mathbb{R}^D$)
40 such that $D \in \mathbb{N}^+$, using image-pair correspondences as input. These dense visual object descriptors
41 provide a generalized representation of objects to a certain extent.

42 The DON framework has shown promise in various domains, including rope manipulation [15], block
43 manipulation [16], robot control [17], fabric manipulation [18], and robot grasp pose estimation [19,
44 20]. Adrian et al. [20] demonstrated that DON can be trained on synthetic data and still generalize
45 well to real-world objects. Furthermore, they highlighted the importance of the dimensionality of the
46 embedding in determining the quality of the descriptors produced by the DON framework.

47 In this paper, we address the challenge posed by the computationally intensive nature of DON and
48 propose a new framework for training DON in a computationally efficient manner. Furthermore, we
49 introduce a novel synthetic data generation pipeline that generates a complete dataset from one image
50 and mask pair. Additionally, the synthetic data generation pipeline does not rely on the noisy depth
51 information produced by today’s consumer-grade depth cameras. We also demonstrate one of the
52 applications of our framework as a robotic grasping pipeline. Our approach aims to contribute to
53 developing a sustainable and efficient, and economical solution for industrial robotics applications.

54 2 Related Work

55 Florence et al. [14] introduced the Pixelwise Contrastive loss function to train DON, which involves
56 sampling pixels in an image-pair and computing the Contrastive loss between the pixels in the first
57 image and those in the second image. This optimization procedure aims to improve the descriptor
58 based on a similarity metric. However, the Pixelwise Contrastive loss function is computationally
59 expensive and requires numerous matching and non-matching image-pair correspondences to work
60 optimally. When optimizing DON using a large number (N) of image-pair correspondences, the
61 computational resources consumed by the optimization procedure increase significantly due to the
62 exponential growth of pixelwise descriptor similarity comparisons (2^N).

63 In their work, Florence [21] discovered that the Pixelwise Contrastive loss function used to train
64 DON might yield poor performance if a computed correspondence is spatially inconsistent. They
65 also highlighted that the precision of contrastive-trained models could be sensitive to the relative
66 weighting between positive and negative sampled pixels. To address these limitations, Florence
67 proposed a new continuous sampling-based loss function called the Pixelwise Distribution loss. This
68 novel loss function leverages smooth and continuous pixel space sampling instead of the discrete pixel
69 space sampling method employed by the Pixelwise Contrastive loss. The Pixelwise Distribution loss
70 eliminates the need for non-matching correspondences, leading to significant savings in computation
71 resources.

72 On a different note, Kupcsik et al. [19] utilized Laplacian Eigenmaps [22] to embed a 3D object model
73 into an optimally generated embedding space, serving as the target for training DON in a supervised
74 fashion. However, this methodology does not reduce the computational resource consumption
75 required to train DON. In contrast, Hadjivelichkov and Kanoulas [23] employed offline unsupervised
76 clustering based on confidence in object similarities to generate hard and soft correspondence labels.
77 These labels were then used as matching and non-matching correspondences to train DON effectively.

78 Building upon the concept of SIMCLR-inspired frameworks [24, 25], Adrian et al. [20] introduced
79 a similar architecture and another novel loss function called the Pixelwise NTXent Loss. This loss
80 function robustly trains DON by leveraging synthetic correspondences computed from image augmen-
81 tations and non-matching image correspondences. Notably, Adrian et al.’s experiments demonstrated
82 that the novel loss function is invariant to batch size variations, unlike the Pixelwise Contrastive Loss.
83 Furthermore, it is worth noting that most of the discussed optimization methodologies heavily rely on
84 correspondences to train DON effectively.

85 Moving on to the aspect of image-pair correspondences and dataset engineering, the DON training
86 strategy proposed in [14, 21] relies on depth information to compute correspondences between image
87 pairs using camera intrinsics and pose information [26]. However, when utilizing consumer-grade
88 depth cameras to capture depth information, the resulting depth data can be noisy, particularly when
89 dealing with tiny, reflecting objects common in industrial environments. Noisy depth information
90 hampers the computation of consistent spatial correspondences in an image pair. To overcome
91 this challenge, Kupcsik et al. [19] associated 3D models of objects with image views, effectively

92 training DON without relying on depth information. Their approach proved efficient for smaller,
93 texture-less, and reflective objects. Additionally, Kupcsik et al. compared different training strategies
94 for producing 6D grasps on industrial objects and demonstrated that a unique supervised training
95 approach enhances pick-and-place resilience in industry-relevant tasks.

96 In contrast, Yen-Chen et al.[27] employed NeRF[28], a method that reconstructs a 3D scene from
97 a sequence of images captured by a smartphone camera. They extracted correspondences from the
98 synthetically reconstructed scene to train DON. Remarkably, Adrian et al.’s experiments indicated
99 that DON trained on synthetic data generalizes well to real-world objects. Furthermore, they adopted
100 the $PCK@k$ metric, commonly used in [29, 30], to evaluate and benchmark DON’s performance in
101 cluttered scenes that were previously not extensively studied.

102 In our work, we do not use any loss functions as proposed in [14, 21, 19, 20, 23, 27] to train DON.
103 However, we adopt the network architecture from DON [14] as our architecture’s backbone and train
104 on the task of the KeypointNet[31, 32] with few network modifications. Moreover, we evaluate the
105 descriptor’s robustness produced by our framework on the $PCK@k$ metric as in comparision to
106 benchmarks in [20] as it is the only benchmark available for DON. Furthermore, we compare the
107 computational resource consumption.

108 3 Methodology

109 In this section, we outline the methodologies employed. Our approach encompasses synthetic dataset
110 engineering, a novel framework, loss function modifications and a comprehensive grasping pipeline.

111 Firstly, we focus on synthetic dataset engineering accomodating spatial, colour and background
112 augmentations. The color and background augmentations help the framework to predict object-oriented
113 descriptors. Secondly, we present a novel framework designed to reduce computational resource
114 consumption with loss function modifications to optimize performance. Lastly, we introduce a
115 comprehensive robot grasping pipeline, integrating all methodologies into a unified system.

116 3.1 Dataset Engineering

117 We have chosen the cap object for creating a synthetic dataset as the cap mesh models are readily
118 available in the Shapenet library [33] as it contains rich object information, including textures.
119 Furthermore, we choose five cap models from the Shapenet library and use Blenderproc [34] to
120 generate the synthetic dataset. We save one RGB image, mask, and depth for each cap model from the
121 synthetic scene. Additionally, we employ synthetic augmentations as proposed in [20] to synthetically
122 spatial augment the cap’s position and rotation in an image, including background randomization using
123 Torchvision [35] library. An augmented image-pair is sampled randomly to generate camera poses for
124 different viewpoints. Additionally, image-pair correspondences are computed¹ as illustrated in the
125 Figure 1. Using depth information, we project the computed correspondences to the camera frame and
126 compute the relative transformation between two camera-frame coordinates of the correspondences
127 using Kabsch’s transformation [36]. Moreover, mask and depth images are not used during inference.

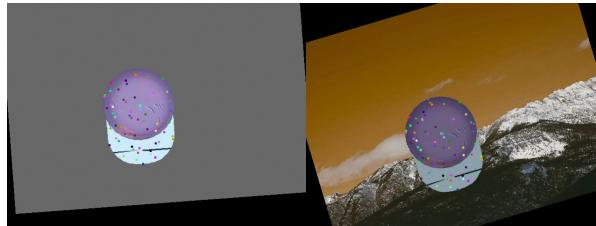


Figure 1: Depiction of image synthetic spatial augmentation and correspondences mapping in an image-pair. The colored encoded dots in the figure represents correspondences in an image-pair.

¹GitHub Link: *link is made anonymous for the review*

128 **3.2 Framework & Mining Strategy**

129 As a backbone, we employ ResNet-34 architecture [37]. We preserve the last convolution layer and
 130 remove the pooling and linear layers. The backbone downsamples the RGB image $I_{RGB} \in \mathbb{R}^{H \times W \times 3}$
 131 to dense features $I_d \in \mathbb{R}^{h \times w \times D}$ such that $h \ll H, w \ll W$ and $D \in \mathbb{N}^+$. We upsample the dense
 132 features from the identity layer (being identical to the last convolution layer in the backbone) as
 133 illustrated in the Figure 2 in page 4 as follows:

$$f_U : I \in \mathbb{R}^{h \times w \times D} \rightarrow I_D \in \mathbb{R}^{H \times W \times D}. \quad (1)$$

134 The upsampled dense features are extracted and treated as dense visual local descriptors produced
 135 from the DON. In otherwords we extract or mine the representations from the backbone. Similarly as
 136 in [31], we stack spatial-probability regressing layer and depth regressing layer on top of the identity
 137 layer to predict $N \in \mathbb{N}^+$ number of keypoint's spatial-probability as follows:

$$f_S : I_d \in \mathbb{R}^{h \times w \times D} \rightarrow I_s^N \in \mathbb{R}^{h \times w \times N}, \quad (2)$$

138 and depth as follows:

$$f_D : I_d \in \mathbb{R}^{h \times w \times D} \rightarrow I_{\hat{d}} \in \mathbb{R}^{h \times w \times N}. \quad (3)$$

139 We incorporate continuous sampling method f_E from [21, 31] to convert the upsampled predicted
 140 spatial-probability and depth of a keypoint to spatial-depth expectation as follows:

$$f_E \circ g_E : [I_s, I_{\hat{d}}]^T \in \mathbb{R}^3, \text{ where } g_E : I \in \mathbb{R}^{h \times w \times N} \rightarrow I \in \mathbb{R}^{H \times W \times N}. \quad (4)$$

141 Furthermore, we train the framework in a twin architecture fashion as proposed in [24, 25, 14, 21, 19,
 142 20, 23, 27] on the modified KeypointNet task.

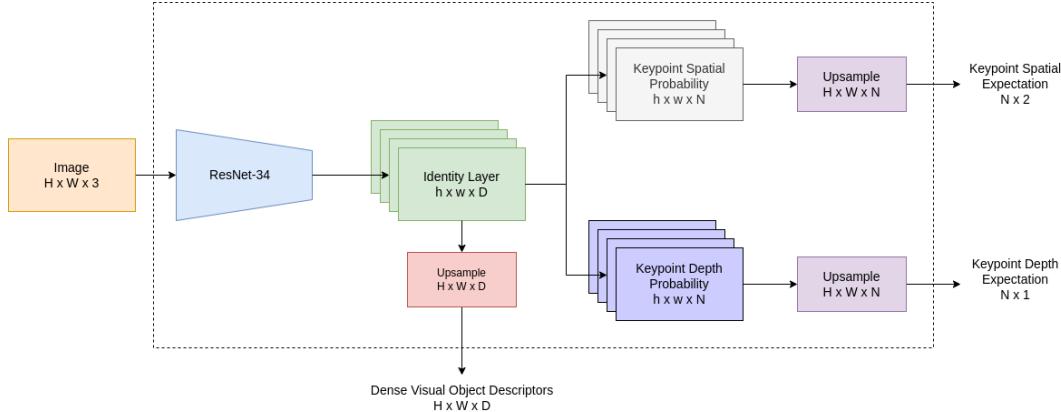


Figure 2: Illustration of the novel framework designed to compute and seamlessly extract dense visual object descriptors efficiently. During inference, we extract dense visual object descriptors directly from the network and ignore predicted spatial-depth expectations of the keypoints.

143 **3.3 Loss Functions**

144 For training, we directly adopt silhouette consistency loss (\mathcal{L}_{obj}), variance loss (\mathcal{L}_{var}) and separation
 145 loss (\mathcal{L}_{sep}) functions from [31] to train the network on the keypoint prediction task. However, we
 146 modify the multi-view consistent loss and relative pose estimation loss. In the case of multi-view
 147 consistency loss we project the predicted spatial-depth expectation using camera intrinsics as follows:

$$X_{cam} \in \mathbb{R}^{3 \times 1} = \mathcal{I}_{cam}^{-1} [u, v, 1.0]^T \times d, \text{ where } \mathcal{I}_{cam} \in \mathbb{R}^{3 \times 3} \text{ and } u, v, d \in \mathbb{R}^+. \quad (5)$$

149 Furthermore, we project the camera coordinates of the keypoints from one camera viewpoint to
 150 another camera viewpoint using relative transformation supplied from the synthetic augmentation
 151 procedure as follows:

$$\mathcal{L}_{mvc} \in \mathbb{R} = \mathcal{H}(\hat{X}_{cam}^B, \mathcal{T}_{A \rightarrow B} \hat{X}_{cam}^A), \text{ where } \hat{X}_{cam} = [X_{cam}, 1.0]^T \in \mathbb{R}^{4 \times 1}, \quad (6)$$

152 In Equation 6, $\mathcal{T}_{A \rightarrow B} \in SE(3)$ is a Special Euclidean Group [38] which is relative transformation
 153 from camera-frame A to camera-frame B . We use Huber loss \mathcal{H} as it produces smoother gradients
 154 for framework optimization. Furthermore, we do not discard the relative transformation information
 155 to calculate the relative pose loss as suggested in [31]. Moreover, being influenced from [32] we
 156 modified the relative pose loss as follows:

$$\mathcal{L}_{pose} = \|log(\mathcal{T}_{truth}^\dagger \mathcal{T}_{pred})\|, \text{ where } log : SE(3) \rightarrow \mathfrak{se}(3) \text{ and } \mathcal{T}^\dagger = \begin{bmatrix} R^T & -R^T t \\ 0^T & 1 \end{bmatrix}. \quad (7)$$

157 3.4 Robot Grasping Pipeline

158 To use the proposed framework as a robot grasping pipeline, we extract dense visual object descriptors
 159 from the network and store one single descriptor of objects in a database manually for now. During
 160 inference, we extract dense visual object descriptors from the network and query the descriptor from
 161 the database to find the closest match as follows:

$$\mathbb{E}[u^*, v^*]_d = \underset{u, v}{\operatorname{argmin}} \exp - \left(\frac{\|I_D[u, v] - d\|}{\exp(t)} \right)^2, \text{ where } \|I_D[u, v] - d\| \in \mathbb{R}^{H \times W}. \quad (8)$$

162 Where $t \in \mathbb{R}$ controls the kernel width influencing the search space to compute the optimal spatial
 163 expectation $\mathbb{E}[u^*, v^*]_d$ of the query descriptor $d \in \mathbb{R}^D$ in the descriptor image $I_D \in \mathbb{R}^{H \times W \times D}$.
 164 The computed spatial expectation is projected to the robot frame using camera intrinsics and poses
 165 to perform a pinch grasp. Furthermore, the Franka Emika 7-DOF robot manipulator with two jaw
 166 gripper and wrist-mounted Intel Realsense D435 camera is used as a testing setup as illustrated in
 167 Figure 3.

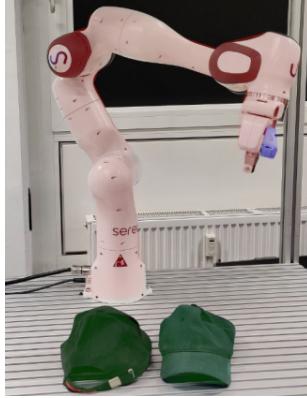


Figure 3: Illustration of the robot grasping pipeline setup. In the image, the robot is highlighted in red, the caps in green, and the camera in blue.

168 4 Experiments & Results

169 In this section, we outline the benchmarking results employed from the methodologies. We benchmark
 170 the DON framework with Pixelwise NT-Xent loss as in [20] and our framework with our revision
 171 of the loss function on a 48GB VRAM GPU. We benchmark the descriptor's robustness with the
 172 $AUC \pm \sigma$ for $PCK@k, \forall k \in [1, 100]$ metric. Furthermore, we benchmark the computational
 173 resource consumption of the DON and our frameworks. We also demonstrate the application of
 174 our framework as a robot-grasping pipeline in two methodologies, one of which our framework
 175 demonstrates its capabilities to produce object-specific 6D poses for robot grasping.

176 4.1 Training Setup

177 We implemented training and benchmarking using “PyTorch-Lightning”[39] and “PyTorch”[40]
 178 libraries. Furthermore, we employ ADAM[41] optimizer to optimize the model for 2500 epochs
 179 with a learning rate of $\alpha = 3 \times 10^{-4}, \beta_1 = 0.9$ and $\beta_2 = 0.999$ with weight decay $\eta = 10^{-4}$ to

180 benchmark the DON with Pixelwise NT-Xent loss as in [20] with a fixed batch size of 1 and 128
181 image-pair correspondences.

182 To train our framework, we employ an ADAM optimizer to optimize the model for 2500 epochs
183 with a learning rate of $\alpha = 1 \times 10^{-3}$, $\beta_1 = 0.9$ and $\beta_2 = 0.999$ with no weight decay. We further
184 use a fixed batch size of 1 and the StepLR scheduler with a step size 2500 and a gamma of 0.9 to
185 train the model with all the loss weights to 1.0 except variance loss weight to 1×10^{-3} . We trained
186 the three models with 128 keypoints with a margin of 2 pixels for each descriptor dimension. We
187 specifically chose 128 keypoints as it aligns with the notion that DON is benchmarked with 128
188 image-pair correspondences.

189 **4.2 Benchmarking & Results**

190 The $AUC \pm \sigma$ for $PCK@k, \forall k \in [1, 100]$ is computed with 256 image-pair correspondences for both
191 models. The metrics mean and std. deviation is calculated from benchmarking three models trained
192 for each descriptor dimension. Due to the limited GPU VRAM capacity, we could not train DON for
193 descriptor dimensions greater than 32. As per Table 1, both frameworks benefit while training them
194 for longer descriptor dimensions. Furthermore, the higher values infer robust descriptors in Table 1.
195 We notice that our framework works robustly as the descriptor dimension gets longer as the metric
196 difference between DON and our framework reduces.

Table 1: Benchmarking outcomes for descriptors' evaluation metric.

Benchmarking for $AUC \pm \sigma$ for $PCK@k, \forall k \in [1, 100]$		
Descriptor Size (D)	Dense Object Nets	Our framework
3	0.922 ± 0.006	0.914 ± 0.009
8	0.933 ± 0.011	0.928 ± 0.015
16	0.948 ± 0.012	0.945 ± 0.010
32	0.953 ± 0.008	0.950 ± 0.009
64	~	0.953 ± 0.006
128	~	0.957 ± 0.012
256	~	0.959 ± 0.008
512	~	0.962 ± 0.011

197 While training both frameworks, we monitor the GPU VRAM consumption. As per the benchmark
198 results in Table 2, the DON consumption increases as the descriptor dimensions get longer while our
199 framework consumes a fraction of the computation resource. Furthermore, lower readings are better
200 in Table 2.

Table 2: Benchmarking outcomes for training computation resource consumption.

Benchmarking for GPU VRAM(GB) consumption		
Descriptor Size (D)	Dense Object Nets	Our framework
3	9.377	4.763
8	13.717	4.785
16	20.479	4.832
32	30.067	4.872
64	~	4.913
128	~	5.409
256	~	6.551
512	~	7.915

201 To check the impact of descriptors' robustness compared to the number of keypoints, we trained our
202 framework with 16 keypoints. Furthermore, we trained three additional models for each descriptor
203 dimension 64, 128, 256, and 512. As per the table 3 compared to the results in Table 1 in page 1,
204 the descriptor's robustness decreased when the framework predicted 16 keypoints. Moreover, this
205 reflects that number of keypoints in our framework and the number of image-pair correspondences in
206 DON are directly proportional to the robustness of the descriptors.

Table 3: Benchmark of our framework with 16 keypoints for GPU VRAM(GB) consumption and $AUC \pm \sigma$ for $PCK@k, \forall k \in [1, 100]$ metric.

Our framework with 16 keypoints		
Descriptor Size (D)	$AUC \pm \sigma$ for $PCK@k, \forall k \in [1, 100]$	VRAM Usage (GB)
64	0.948 ± 0.009	3.799
128	0.952 ± 0.010	4.191
256	0.955 ± 0.013	5.241
512	0.957 ± 0.006	7.341

207 4.3 Descriptor Inspection

208 Furthermore, to inspect the results of trained DON, an interface is built using the PyGame library [42]
209 to visualize the results of the trained DON. The mouse pointer in the image space is mapped to the
210 pixel, and the descriptor at that pixel is queried in another image-descriptor space. We further use
211 the spatial probability of the descriptor to visualize the queried descriptor in the image space using
212 Equation 8 in page 5. We identify if there are any multi-modal spatial activations in the descriptor
213 spaces and none, as shown in Figure 4.

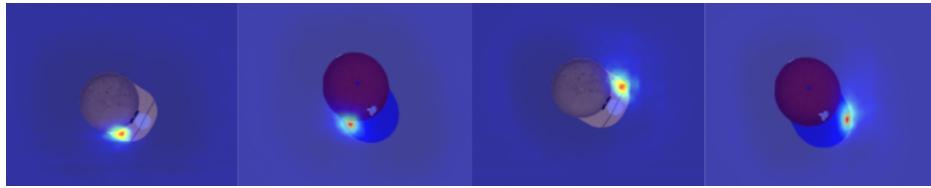


Figure 4: Depiction of the spatial probability heatmaps of the descriptor in the image space. We set the temperature in the Equation 8 to 1.1 and render the spatial probability heatmaps in the interface. The first and second image from the left and the right highlights the semantically equivalent descriptors in the image space.

214 For the robot grasping pipeline, we trained our framework with actual caps. As the synthetic data
215 generation only needs mask and depth information, we could create a mask in no time. Additionally,
216 while training the framework, we do not need the actual real-world depth information as it computes
217 its own. We later extracted the dense visual local descriptors from the framework. We visually
218 inspected for any inconsistencies in the descriptor space, as shown in Figure 5, and found it consistent.
219 Furthermore, we did not use the models trained on the synthetic dataset, as the representations were
220 inconsistent with the real caps.



Figure 5: The image depicts the visual inspection of the dense visual descriptors space of the real caps using our developed interface. We trained our framework on the first two caps from the left, and our framework could generalize the object representations on an unseen cap while training illustrated in the first image from the right.

221 **4.4 Robot Grasping Pipeline**

222 For robot grasping, a descriptor is picked from the descriptor space and queried in real-time such that
 223 robot can pinch-grasp the object. We could successfully grasp the caps with the robot, as shown in
 224 Figure 6 in page 8. Furthermore, we did not evaluation the robot grasping itself.

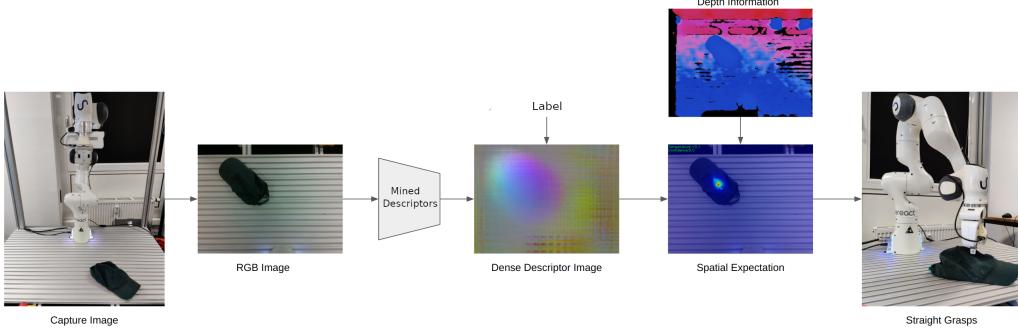


Figure 6: Depiction of the straight robot grasping pipeline.

225 As our framework inertly regresses keypoints on the object, we could use it as an alternative approach
 226 to grasp the caps by computing the pose generated by the keypoints considering the actual depth
 227 information instead of network-regressed depth information. We extract the spatial probability of
 228 each keypoint from the framework and deactivate spatial probabilities where the depth information is
 229 missing, as the depth image from the camera is noisy. Furthermore, the spatial expectations of the
 230 keypoints are projected to the camera frame to calculate a 6D pose in the camera frame. The 6D pose
 231 is transformed in the robot frame to perform an aligned grasp, as shown in Figure 7.

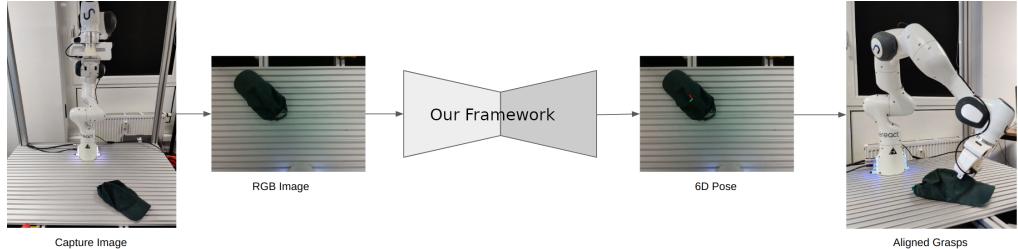


Figure 7: Illustration of the aligned robot grasping pipeline.

232 **5 Conclusion**

233 This paper introduces a novel framework for mining dense visual object descriptors without explicitly
 234 training DON. We have successfully eliminated the requirement for image-pair correspondence
 235 mapping in training DON by employing synthetic augmentation data generation and a novel deep-
 236 learning architecture. Our benchmarking results showcase the effectiveness of our framework in
 237 generating robust and denser visual local descriptors. However, it needs to outperform the original
 238 DON framework in robustness. Moreover, a notable advantage of our proposed framework is its
 239 significantly reduced computational resource consumption, amounting to a remarkable 86.67%
 240 decrease compared to the originally proposed framework. It is important to note that our current
 241 framework is limited to single object-dense visual descriptors. Nevertheless, we have plans to
 242 extend our methodology to encompass the production of multi-object dense visual descriptors in
 243 cluttered scenes. By doing so, we aim to enhance the versatility and applicability of our framework in
 244 real-world scenarios. To demonstrate the practicality of our framework, we have integrated it into a
 245 robot-grasping pipeline using two distinct methodologies. Remarkably, our framework can generate
 246 object-specific 6D poses, enhancing robot grasping performance. This successful application further
 247 highlights the potential utility of our framework in real-world robotic systems.

248 **References**

- 249 [1] N. Blomkamp, H. Zimmer, and S. Kinberg. *Chappie*. Sony Pictures Home Entertainment, 2015.
- 250 [2] G. Lucas and S. Ulstein. *Star wars*. 20th Century-Fox, 1977.
- 251 [3] D. Silver, T. Hubert, J. Schrittwieser, I. Antonoglou, M. Lai, A. Guez, M. Lanctot, L. Sifre, D. Kumaran, T. Graepel, et al. “A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play”. In: *Science* 362.6419 (2018), pp. 1140–1144.
- 254 [4] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, et al. “Mastering the game of Go with deep neural networks and tree search”. In: *nature* 529.7587 (2016), pp. 484–489.
- 257 [5] B. Entertainment. *World of warcraft*. Insight Editions, Division Of Palac, 2013.
- 258 [6] A. Krizhevsky, I. Sutskever, and G. E. Hinton. “Imagenet classification with deep convolutional neural networks”. In: *Communications of the ACM* 60.6 (2017), pp. 84–90.
- 260 [7] J. Long, E. Shelhamer, and T. Darrell. “Fully convolutional networks for semantic segmentation”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2015, pp. 3431–3440.
- 262 [8] K. He, G. Gkioxari, P. Dollár, and R. Girshick. “Mask r-cnn”. In: *Proceedings of the IEEE international conference on computer vision*. 2017, pp. 2961–2969.
- 264 [9] R. A. Güler, N. Neverova, and I. Kokkinos. “Densepose: Dense human pose estimation in the wild”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2018, pp. 7297–7306.
- 266 [10] OpenAI. *GPT-4 Technical Report*. 2023. arXiv: 2303.08774 [cs.CL].
- 267 [11] A. Chowdhery et al. “PaLM: Scaling Language Modeling with Pathways”. In: *arxiv:2204.02311* (2022).
- 268 [12] E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell. “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” In: *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 2021, pp. 610–623.
- 271 [13] E. Strubell, A. Ganesh, and A. McCallum. “Energy and policy considerations for deep learning in NLP”. In: *arXiv preprint arXiv:1906.02243* (2019).
- 273 [14] P. R. Florence, L. Manuelli, and R. Tedrake. “Dense object nets: Learning dense visual object descriptors by and for robotic manipulation”. In: *arXiv preprint arXiv:1806.08756* (2018).
- 275 [15] P. Sundaresan, J. Grannen, B. Thananjeyan, A. Balakrishna, M. Laskey, K. Stone, J. E. Gonzalez, and K. Goldberg. “Learning Rope Manipulation Policies Using Dense Object Descriptors Trained on Synthetic Depth Data”. In: *CoRR* abs/2003.01835 (2020). arXiv: 2003.01835.
- 278 [16] C.-Y. Chai, K.-F. Hsu, and S.-L. Tsao. “Multi-step Pick-and-Place Tasks Using Object-centric Dense Correspondences”. In: *2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2019, pp. 4004–4011. DOI: 10.1109/IROS40897.2019.8968294.
- 281 [17] P. Florence, L. Manuelli, and R. Tedrake. “Self-supervised correspondence in visuomotor policy learning”. In: *IEEE Robotics and Automation Letters* 5.2 (2019), pp. 492–499.
- 283 [18] A. Ganapathi et al. “Learning Dense Visual Correspondences in Simulation to Smooth and Fold Real Fabrics”. In: *2021 IEEE International Conference on Robotics and Automation (ICRA)*. 2021, pp. 11515–11522. DOI: 10.1109/ICRA48506.2021.9561980.
- 286 [19] A. Kupcsik, M. Spies, A. Klein, M. Todescato, N. Waniek, P. Schillinger, and M. Bürger. “Supervised Training of Dense Object Nets using Optimal Descriptors for Industrial Robotic Applications”. In: *arXiv preprint arXiv:2102.08096* (2021).
- 289 [20] D. B. Adrian, A. G. Kupcsik, M. Spies, and H. Neumann. “Efficient and Robust Training of Dense Object Nets for Multi-Object Robot Manipulation”. In: *2022 International Conference on Robotics and Automation (ICRA)*. IEEE. 2022, pp. 1562–1568.
- 292 [21] P. R. Florence. “Dense visual learning for robot manipulation”. PhD thesis. Massachusetts Institute of Technology, 2020.
- 294 [22] M. Belkin and P. Niyogi. “Laplacian eigenmaps for dimensionality reduction and data representation”. In: *Neural computation* 15.6 (2003), pp. 1373–1396.
- 296 [23] D. Hadjivelichkov and D. Kanoulas. “Fully Self-Supervised Class Awareness in Dense Object Descriptors”. In: *5th Annual Conference on Robot Learning*. 2021.
- 298 [24] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton. “A simple framework for contrastive learning of visual representations”. In: *International conference on machine learning*. PMLR. 2020, pp. 1597–1607.
- 300 [25] J. Zbontar, L. Jing, I. Misra, Y. LeCun, and S. Deny. “Barlow twins: Self-supervised learning via redundancy reduction”. In: *International Conference on Machine Learning*. PMLR. 2021, pp. 12310–12320.
- 303 [26] R. Hartley and A. Zisserman. *Multiple view geometry in computer vision*. Cambridge university press, 2003.
- 305 [27] L. Yen-Chen, P. Florence, J. T. Barron, T.-Y. Lin, A. Rodriguez, and P. Isola. *NeRF-Supervision: Learning Dense Object Descriptors from Neural Radiance Fields*. 2022. DOI: 10.48550/ARXIV.2203.01913.

- 307 [28] B. Mildenhall, P. P. Srinivasan, M. Tancik, J. T. Barron, R. Ramamoorthi, and R. Ng. “Nerf: Representing
 308 scenes as neural radiance fields for view synthesis”. In: *Communications of the ACM* 65.1 (2021), pp. 99–
 309 106.
- 310 [29] C.-Y. Chai, K.-F. Hsu, and S.-L. Tsao. “Multi-step pick-and-place tasks using object-centric dense
 311 correspondences”. In: *2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*.
 312 IEEE. 2019, pp. 4004–4011.
- 313 [30] M. E. Fathy, Q.-H. Tran, M. Z. Zia, P. Vernaza, and M. Chandraker. “Hierarchical metric learning and
 314 matching for 2d and 3d geometric correspondences”. In: *Proceedings of the european conference on*
 315 *computer vision (ECCV)*. 2018, pp. 803–819.
- 316 [31] S. Suwajanakorn, N. Snavely, J. J. Tompson, and M. Norouzi. “Discovery of latent 3d keypoints via
 317 end-to-end geometric reasoning”. In: *Advances in neural information processing systems* 31 (2018).
- 318 [32] W. Zhao, S. Zhang, Z. Guan, W. Zhao, J. Peng, and J. Fan. “Learning deep network for detecting 3d
 319 object keypoints and 6d poses”. In: *Proceedings of the IEEE/CVF Conference on computer vision and*
 320 *pattern recognition*. 2020, pp. 14134–14142.
- 321 [33] A. X. Chang, T. Funkhouser, L. Guibas, P. Hanrahan, Q. Huang, Z. Li, S. Savarese, M. Savva, S. Song,
 322 H. Su, et al. “Shapenet: An information-rich 3d model repository”. In: *arXiv preprint arXiv:1512.03012*
 323 (2015).
- 324 [34] M. Denninger, M. Sundermeyer, D. Winkelbauer, Y. Zidan, D. Olefir, M. Elbadrawy, A. Lodhi, and
 325 H. Katam. “Blenderproc”. In: *arXiv preprint arXiv:1911.01911* (2019).
- 326 [35] S. Marcel and Y. Rodriguez. “Torchvision the machine-vision package of torch”. In: *Proceedings of the*
 327 *18th ACM international conference on Multimedia*. 2010, pp. 1485–1488.
- 328 [36] W. Kabsch. “A solution for the best rotation to relate two sets of vectors”. In: *Acta Crystallographica*
 329 *Section A: Crystal Physics, Diffraction, Theoretical and General Crystallography* 32.5 (1976), pp. 922–
 330 923.
- 331 [37] K. He, X. Zhang, S. Ren, and J. Sun. “Deep residual learning for image recognition”. In: (2016), pp. 770–
 332 778.
- 333 [38] W. P. Thurston. “Three-Dimensional Geometry and Topology, Volume 1”. In: *Three-Dimensional Geome-*
 334 *try and Topology, Volume 1*. Princeton university press, 2014.
- 335 [39] W. A. Falcon. “Pytorch lightning”. In: *GitHub* 3 (2019).
- 336 [40] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein,
 337 L. Antiga, et al. “Pytorch: An imperative style, high-performance deep learning library”. In: *Advances in*
 338 *neural information processing systems* 32 (2019).
- 339 [41] D. P. Kingma and J. Ba. “Adam: A method for stochastic optimization”. In: *arXiv preprint*
 340 *arXiv:1412.6980* (2014).
- 341 [42] S. Kelly. “Basic introduction to pygame”. In: *Python, PyGame and Raspberry Pi Game Development*.
 Springer, 2016, pp. 59–65.